

# AMAZON DYNAMODB



## A FULLY-MANAGED DATABASE WITH HIGH PERFORMANCE AT ANY SCALE

### INTRODUCTION 🌐

Why should you care about DynamoDB?  
Because it's **fully managed**, **highly available** & **scales on-demand** with **low latencies**.  
For getting you hooked, at Prime Days 2021 DynamoDB served **89.2 million requests per second** at its peak!

### HOT PARTITIONS 🔥

Your provisioned read & write capacity units will be distributed among all internal partitions.  
If your partition keys are not well-distributed, it will be easier to get your requests throttled as a subset of your partitions (or worse: a single) can receive the **majority** of read and/or write requests.

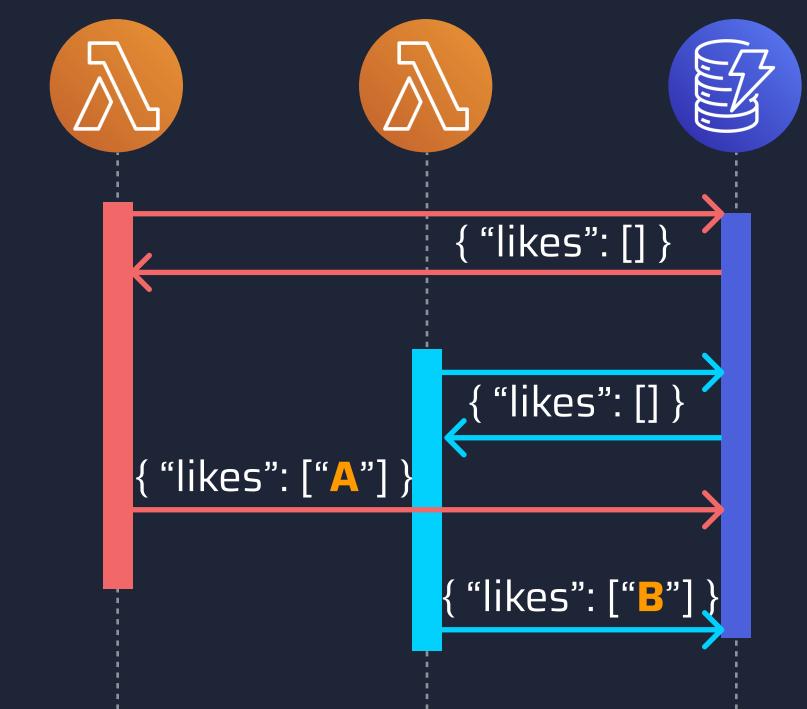


### RACE CONDITIONS ⚡

Often, there are possible race conditions due to multi-tenancy where writes can be **lost**.

Example:

1. **Lambda A** reads Document X
2. **Lambda B** reads Document X
3. **Lambda A** writes Document X
4. **Lambda B** writes Document X ⚡



### UPDATE EXPRESSIONS ◆

Safely modify certain fields of a complex document.

- **SET** - adding one or several attributes to an item.
- **REMOVE** - removing one or several attributes.
- **ADD** - adding new attributes with its value.
- **DELETE** - remove one or more elements from a set.

### CAPACITY MODES 📈

You can choose between those two types, which can be changed **any time**.

- **provisioned** - specifying the capacity units for your table & you'll be independently charged of your actual usage.
  - **on-demand** - you'll be charged per request.
- 💡 The best fit highly depends on your **traffic patterns**.

### DATA TYPES ◆

Besides your primary key, your document can contain other fields of different **Scalar**, **Document** and **Set Types**.

#### SCALAR

Represented by exactly one value.

- String (**S**)
- Number (**N**)
- Binary (**B**)
- Boolean (**BOOL**)
- Null (**NULL**)

#### DOCUMENT

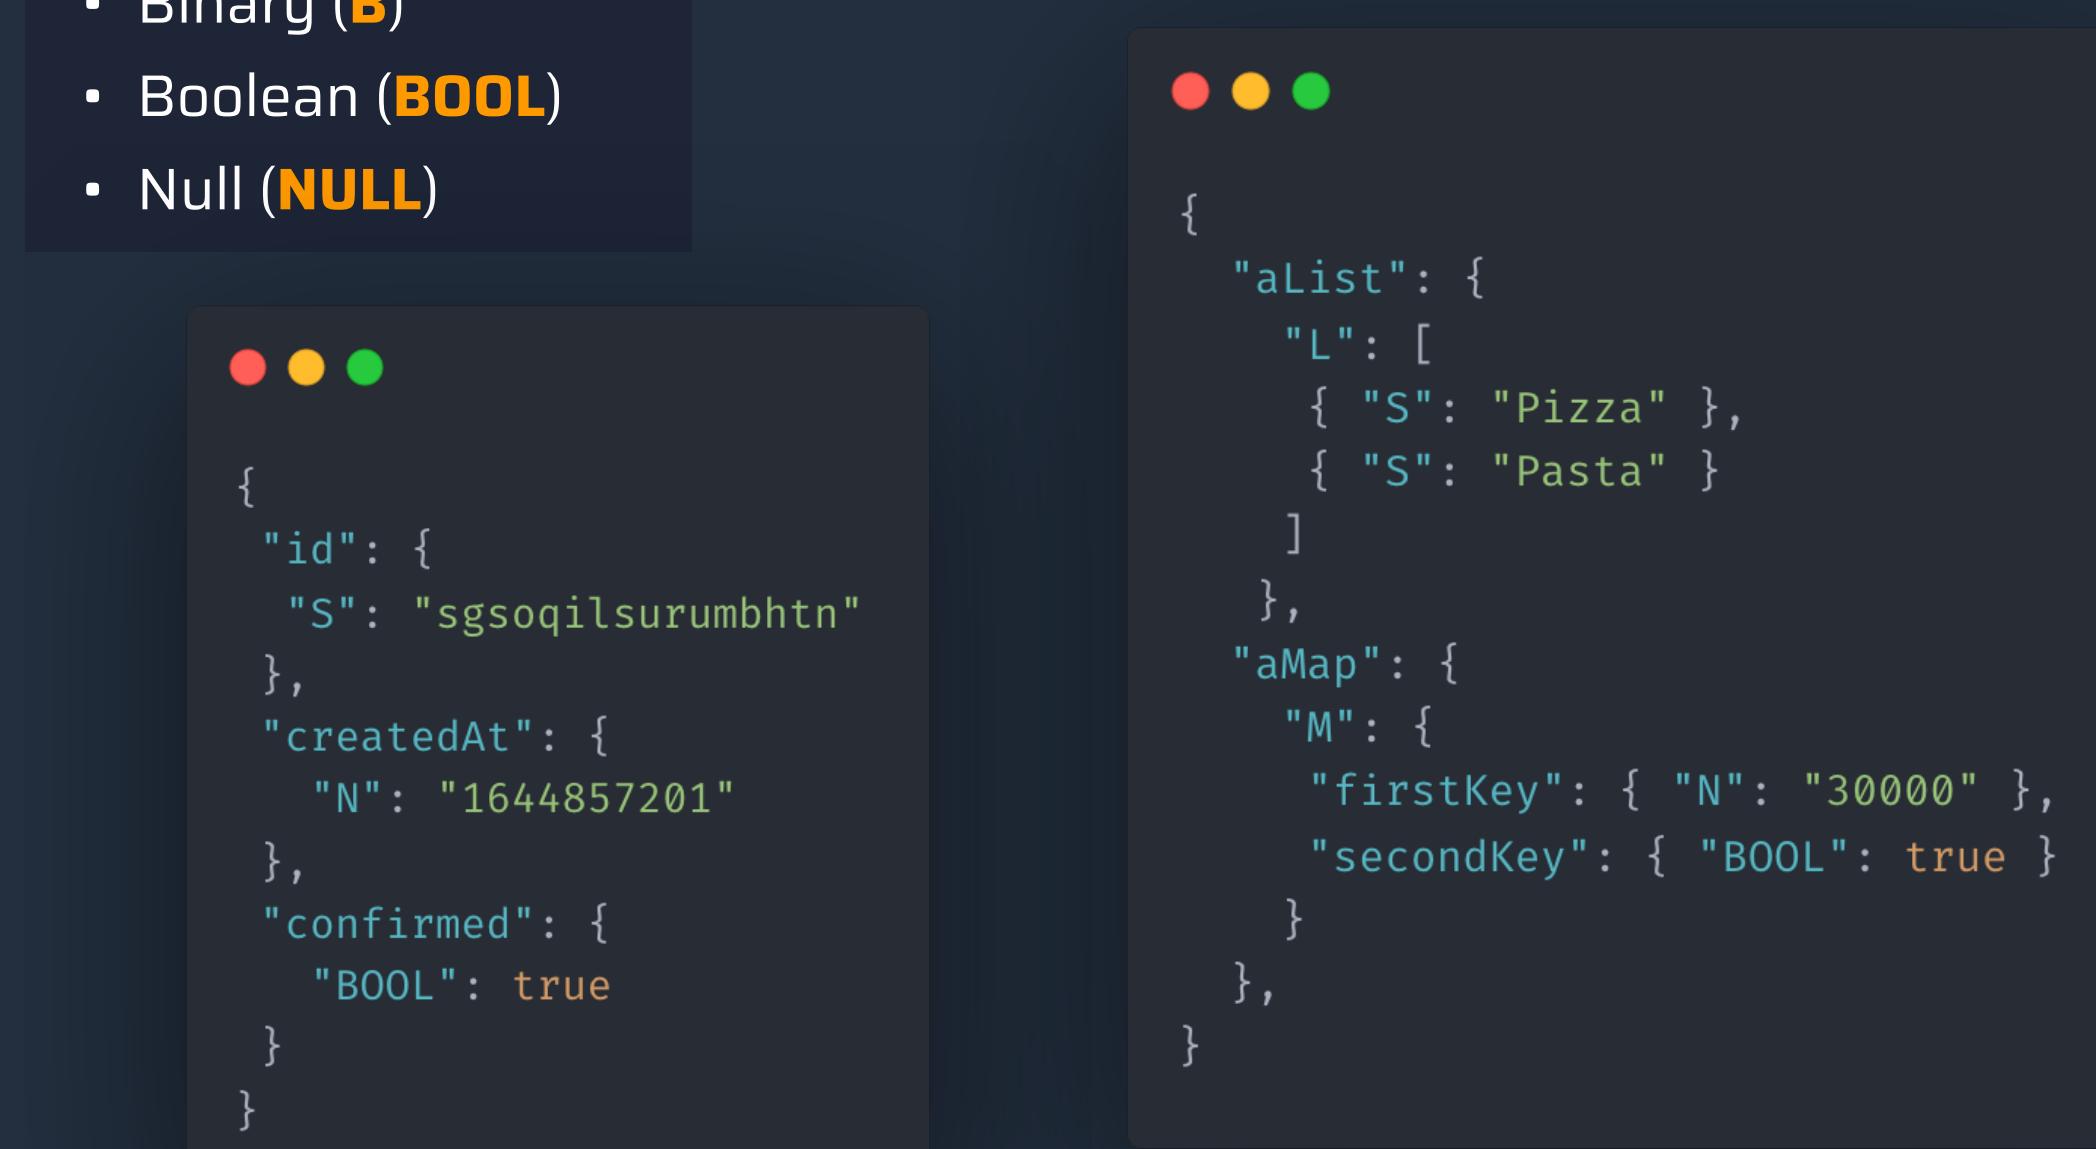
Complex structure with nested attributes.

- List (**L**)
- Map (**M**)

#### SET

Multiple scalar values.

- String Set (**SS**)
- Number Set (**NS**)
- Binary Set (**BS**)



### OPTIMISTIC LOCKING 🔒

DynamoDB's **Optimistic Locking** allows us to verify that we're really updating the item we're expecting by using a dedicated version field.

In the previous example, both read operations would read **version 1** and also expect to write **version 2**. The first operation would succeed, but the second would fail as the expected version would not match with the expectations.



### ON-DEMAND VS. PROVISIONED 💸

Go with **On-Demand** if you got unpredictable traffic, as it scales immediately and you're only paying for what you actually use.

With steady load or known patterns, pick **provisioned** as it can be up to almost **7 times** less expensive.

Rule of thumb:

- variable, **unpredictable** traffic → **on-demand**
- variable, **predictable** traffic → **provisioned**
- **steady, predictable** traffic → **reserved**

### RETRIEVING ITEMS ↗

That's where it gets interesting and you see differences to SQL or other NoSQL solutions.

You can only **query** via your partition key (and sort key condition, if there's any) of your main or secondary indexes.

Everything else needs **scan**.

### SCAN 🌊

A scan is just running through your table looking for items that are matching your expression.

You'll be **charged for the items that are scanned**, not the items that are retrieved.

### OBSERVABILITY 🔍

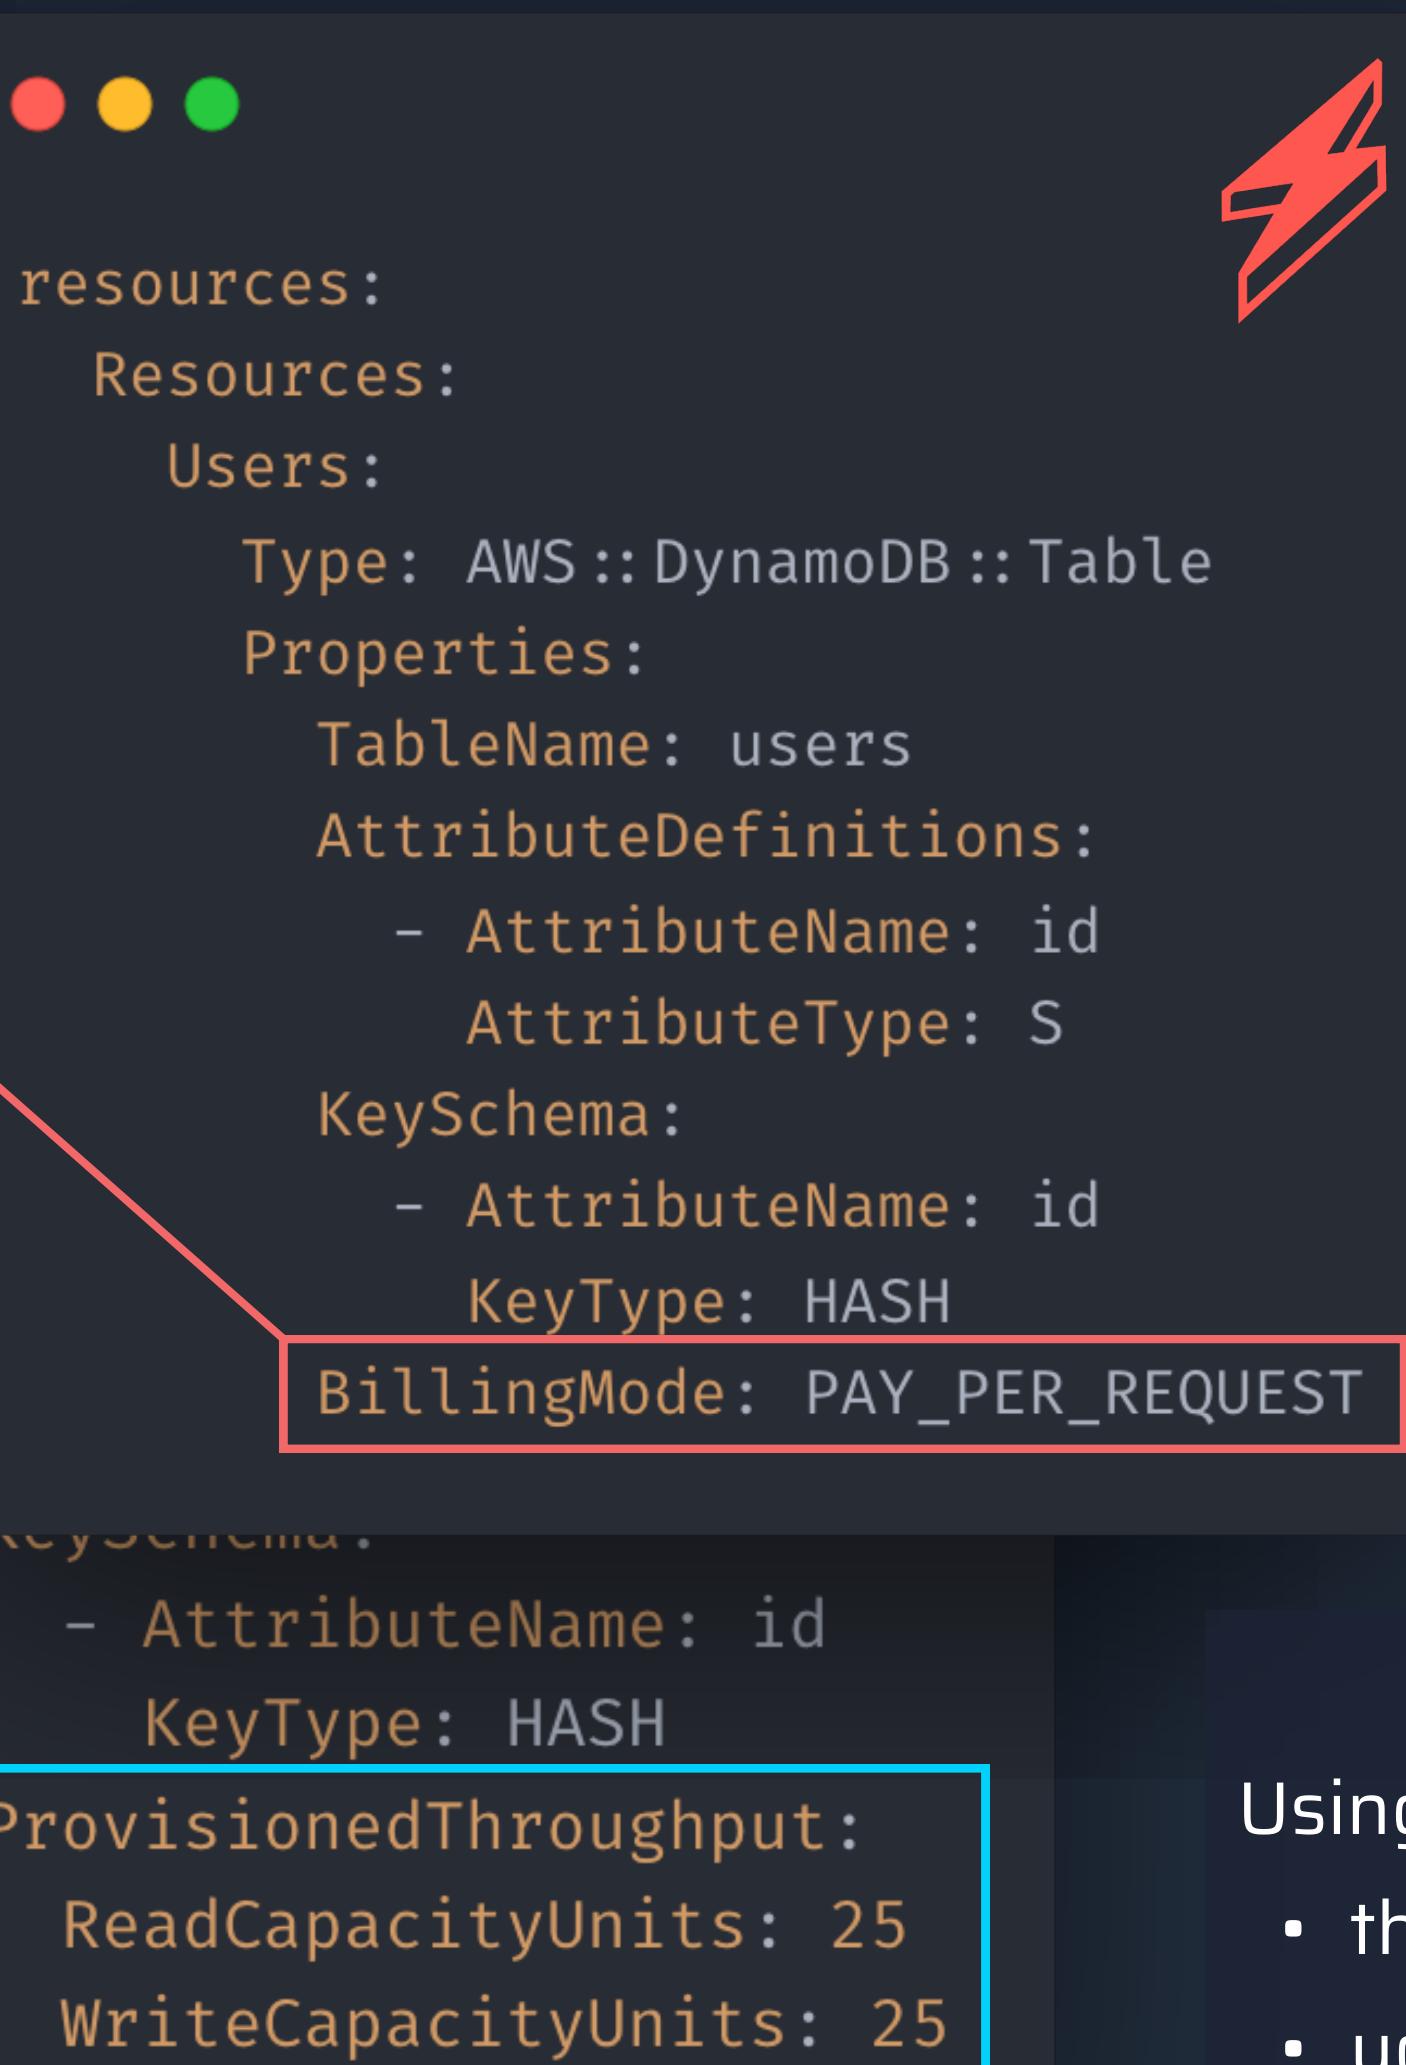
CloudWatch comes with default metrics like **used read & write capacity units** or the **number of throttles**. It helps to analyse DynamoDB traffic patterns so you can optimise for low costs.

**Third-party tools** like **Dashbird.io** help you with **Slack notifications** for critical events like throttles & give you general guidance with **well-architected tips**.

DYNAMODB - 41 resources info operational  
DynamoDB does not have continuous backups enabled

### FREE TIER 💰

Your first year includes **25 Read & Write Capacity Units** each for **provisioned** capacity mode every month.



### BASIC CONCEPTS 📄

In comparison to SQL, a document in DynamoDB doesn't have a fixed schema. What is defined by the table: the **primary key**, which **uniquely** identifies a single document.

### PRIMARY KEYS 🏆

It's a documents **unique** identifier & must be provided when inserting a new item.

There are two different types of primary keys:

- **simple** - your partition key & therefore a single field.
- **composite** - a combination of your partition & sort key.

### PARTITION KEYS 🖐️

Internally, DynamoDB consists of **different partitions** and your partition key will determine the partition where an item will be stored at.



### SORT KEYS 🎯

Using sort keys come in with benefits:

- they can be used with operators in queries such as **begins\_with**, **between**, **>**, **<**, and so on.
- you can create hierarchical (1:n) relationships - citing an example from the docs:  
**[country]#[region]#[state]#[county]#[city]#[neighborhood]**

💡 The combination of a partition & sort key (= your primary key) needs to be **unique**.

### BACKUPS 🔒

AWS offers you a lot of options for backing up your tables.

- **On-Demand Backups** - trigger backups manually or via a scheduled event.
- **Continuous Backups** via **Point-In-Time-Recovery (PITR)** - enables you to restore your table to any state with the last 35 days
- **Exporting Backups to S3** - export your data to an S3 bucket. You can also automatically trigger this process via Lambda & EventBridge rules.

### SECONDARY INDEXES 🔑

You can create two types of indexes, which are specifying **alternative key structures**. Those can also be used to query items.

- **Local Secondary Index (LSI)** - needs to have the same partition key, but an alternative sort key.
- **Global Secondary Index (GSI)** - partition & sort key can be different.

### STREAMS 💬

Asynchronously trigger invocations of other services like Lambda when an item is **created**, **updated** or **deleted**.



aws